

Project Details

- Author: Felipe Paiva Alencar
- Subject: Embedded Systems - Pascal Benoit
- University: Polytech Montpellier
- Date: June 2023

Features

• RISC-V Core:

- DIY implementation of the RV32I instruction set architecture (ISA), providing a flexible and customizable processing core.
- Uses the AXI4-Lite memory interface enabling a seamless interface with different memory devices.
- Has support for interrupts allowing the implementation of event-driven functionalities.

• Graphics:

- The sprite architecture presented in class, has been enhanced to enable dynamic changes to sprite contents, multiple texture scales, and memory sharing between sprites.

- Fully parametric, has support 80 sprites arranged in 5 clusters in the default configuration.

• Memory:

- 4 kilobytes of ROM: Storing a small bootloader that also performs a quick test of the ISA implementation.
- 64 kilobytes of RAM: More than sufficient memory capacity to handle game, texture, and world data.
- 20.48 kilopixels of texture memory: arranged in 5 clusters of sprites.

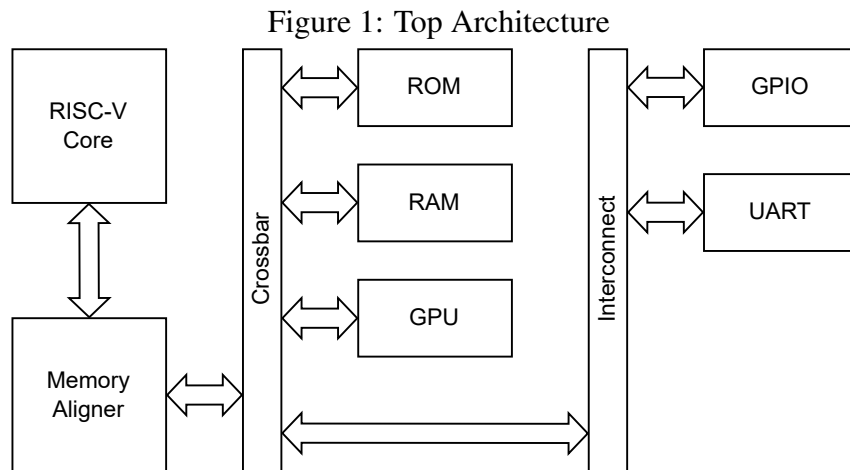
• Peripherals:

- General Purpose Input/Output (GPIO) that allows the interface of the software with the buttons and switches.
- Universal asynchronous receiver-transmitter (UART) that provides a communication channel.

• Build Tools:

- The build and flashing process is efficiently automated by a well designed Makefile, simplifying the compilation, texture packaging, world generation, and linking tasks.
- Support for building and flashing via a single command, saving valuable development time and effort.

Hardware



The top architecture is composed of a RISC-V core, a memory aligner, a crossbar, read-only memory (ROM), random-access memory (RAM) and simple 2D graphics accelerator. Furthermore, the GPIO and UART are classified as peripherals and communicate with the CPU via a simpler interconnect.

Everything is interconnected with the help of AXI4-Lite interfaces. The crossbar and interconnect are imported from the `verilog-axi` library. Everything else, including the RISC-V core, was completely developed in-house. Both of them are overkill for this context, but this architecture was conceived with the use of multiple cores in mind where the use of crossbars and interconnects would be fundamental.

The RISC-V core implements the `rv32i` Base Integer Instruction Set. The `Zicsr` Control and Status Register (CSR) Instructions were also implemented because they are essential for the implementation of interrupts. Interrupts were implemented, but are not currently used for anything useful. The core is multicycle but not pipelined. A simple ISA test is performed by the bootloader in order to maintain one's mental well-being during the process of software development.

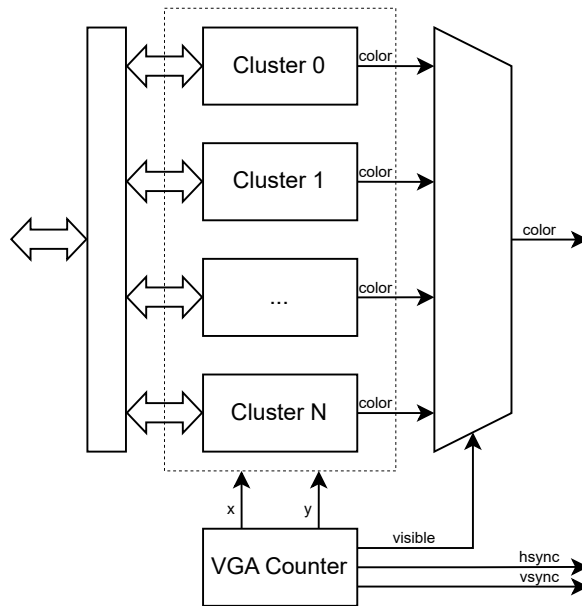
The read-only memory (ROM) is 4 kilobytes long and is used by the bootloader. The use of a bootloader was developed to avoid having to rebuild the FPGA bitstream every software iteration.

The random-access memory (RAM) is 64 kilobytes long and is used for the C code, texture data and world data. The texture data is initially located in RAM, but is moved to the video memory before the start of the game. This allows for changes in the textures without having to rebuild the bitstream. In fact, one could implement a completely different game without rebuilding the hardware implementation.

The General Purpose Input/Output (GPIO) implementation is very basic consisting of simple registers. Therefore, there is no hardware support for changing the GPIO state with a mask.

The Universal Asynchronous Receiver Transmitter (UART) operates at 115200 bauds with 8 data bits, 1 stop bit and no parity bits. This interface is fundamental for the bootloader and is very helpful during software development since it provides some kind of debugging support through the use of the `printf` function.

Figure 2: GPU Architecture

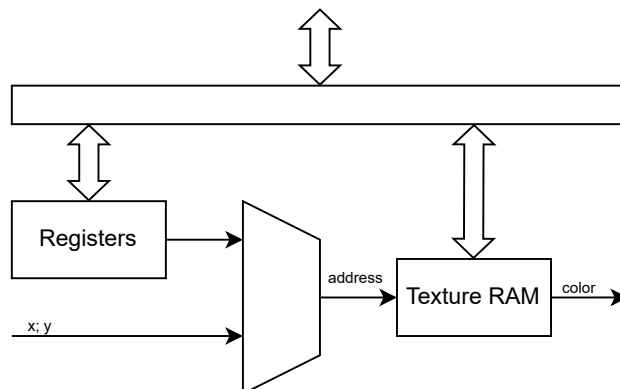


The graphics accelerator developed in class was greatly overhauled into the architecture described in the schema above. The renewed architecture is fully parametric and is capable of displaying 80 sprites with transparency support in the current configuration. These sprites are divided in 5 clusters. A cluster is a group of sprites that uses the same texture RAM. This division was introduced to mitigate the hardware size. A perk of this architecture is that transparency does not work between overlapping sprites in the same cluster. Furthermore, a sprite is not obliged to use the whole texture. A sub-region can be configured for each sprite through register access in software.

The GPU has AXI4-Lite interface that allows for cluster register access. This interface is not propagated to the clusters because its features are not needed and using a simpler interface reduces hardware complexity. Therefore, a converter is present in the GPU top architecture. The converter is not present in the schema.

The VGA Counter provides the VGA synchrony signals and the current position of the electron beam in the screen. This position is used by each cluster to compute the pixel color value. Finally, combinatory logic decides the final value of the VGA color signals based upon the color value of each cluster, the visibility signal of the VGA Counter and the cluster priority order. Transparency is implemented resignifying the value 0xFF to mean transparency. The sky color is hardcoded in this logic.

Figure 3: Cluster Architecture



Each cluster is composed of registers that store the position and texture coordinates of each sprite. With the position value, combinatory logic computes the address that stores the current color value. Then, this value is retrieved in the texture RAM. A very simple interconnect provides a write-only interface for the registers and RAM.

Software

In this case, a final project of the embedded systems subject, I believe that the tools used to build the software (the game) are more relevant than the game itself. Therefore, much thought and time were put into these tools.

I choose the Clang/LLVM toolchain due to its inherent focus on extensibility. By utilizing this toolchain, I could, for example, seamlessly integrate my own instructions into the assembler. However, it's important to note that no custom instructions have been implemented thus far.

The `make` tool is extensively used to automate the building process, determining the source code files that require recompilation and specifying how they should be compiled. It accomplishes this by leveraging a set of rules and dependency relationships defined in the `Makefile`. However, as an experienced programmer, I have encountered more flawed examples of `Makefiles` than commendable ones. Despite its imperfections, I have invested considerable thought and consideration into this project's `Makefile`. Additionally, I have integrated a few helpful scripts into the build process. For instance, you can effortlessly execute `make flash` to compile all textures, generate the world, compile the code, link everything, and flash the resulting binary.

Moving on to the code itself, let's delve into how the interface with the hardware is implemented. It's important to emphasize that, in this architecture, the specialized is essentially just a region of memory for the CPU. To handle this, the `hardware.h` header contains a set of definitions that abstract the memory mapping.

Among the `#defines`, you will find `structs` that provide information regarding how data is organized in the different hardware registers. Additionally, there are further `#defines` that enable the conversion of pure addresses into pointers to these specialized structures.

The usage of the `volatile` qualifier is crucial in this context. It signals to the compiler that the associated variables may undergo unexpected changes. By doing so, it prevents unwanted optimizations that could potentially result in incorrect program behavior.

Ultimately, a robust and intuitive interface to the hardware is achieved, which greatly facilitates software development and minimizes the risk of bugs.

Following the previous section, we come across the `mylib.c` file, where we have provided implementations for functions typically offered by the standard library. I opted out of using the standard library, and as a result, these functionalities are implemented on an as-needed basis. A significant portion of this code is dedicated to the implementation of the `printf` function, which proves to be immensely useful in the debugging process. It's important to note that this function was not developed by me. If you are interested in it, please refer to the references.

Figure 4: World Render Flowchart

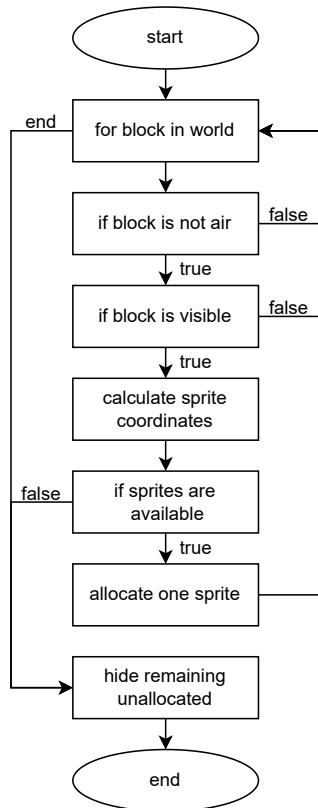
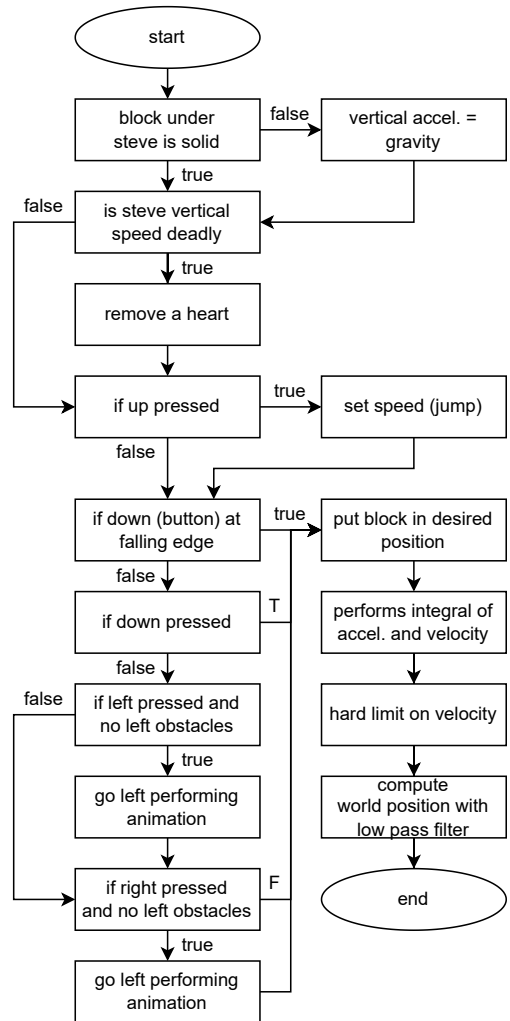


Figure 5: Physics Flowchart



I won't go into extensive explanations about the game mechanics algorithms as I don't consider them as the main focus of the project. If the reader seeks a detailed description of the implementation, they can refer to the source code itself, which serves as the ultimate reference. However, I will present a simplified overview of the key components that merit attention: the world "rendering" algorithm, which efficiently handles 8192 blocks using only 64 block sprites, and the game physics algorithm, which, although flawed and convoluted, can be better understood with the aid of a flowchart for readers interested in exploring the code further.